

Class 9: Candy Mini-Project

Joseph Lo (PID: A18121493)

Table of contents

Background	1
Data Import	1
Exploratory analysis	4
Overall Candy Rankings	7
Taking a look at pricepercent	13
Exploring the correlation structure	14
Principal Component Analysis (PCA)	15
Summary	17

Background

In this mini-project, you will explore FiveThirtyEight's Halloween Candy dataset.

We will use lots of **ggplot** some basic stats, correlation analysis and PCA to make sense of the landscape of US candy - something hopefully more relatable than the proteomics and transcriptomics work that we will use these methods on throughout the rest of the course.

Data Import

Our dataset is a CSV file so we use `read.csv()`

```
candy <- read.csv("candy-data.txt", row.names=1)
head(candy)
```

	chocolate	fruity	caramel	peanut	yalmond	nougat	crisped	rice	wafer
100 Grand	1	0	1		0	0			1
3 Musketeers	1	0	0		0	1			0
One dime	0	0	0		0	0			0

	0	0	0	0	0	0
One quarter	0	0	0	0	0	0
Air Heads	0	1	0	0	0	0
Almond Joy	1	0	0	1	0	0
	hard	bar	pluribus	sugarpercent	pricepercent	winpercent
100 Grand	0	1	0	0.732	0.860	66.97173
3 Musketeers	0	1	0	0.604	0.511	67.60294
One dime	0	0	0	0.011	0.116	32.26109
One quarter	0	0	0	0.011	0.511	46.11650
Air Heads	0	0	0	0.906	0.511	52.34146
Almond Joy	0	1	0	0.465	0.767	50.34755

Q1. How many different candy types are in this dataset?

```
nrow(candy)
```

[1] 85

There are 85 in this dataset

Q2. How many fruity candy types are in the dataset?

```
sum(candy$fruity)
```

[1] 38

38 fruity candy types.

Q3. What is your favorite candy (other than Twix) in the dataset and what is it's winpercent value?

```
candy["Milky Way", "winpercent"]
```

[1] 73.09956

```
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

`intersect`, `setdiff`, `setequal`, `union`

```
candy |> filter(row.names(candy) == "Milky Way") |>
  select(winpercent)
```

Milky Way 73.09956

Q4. What is the winpercent value for “Kit Kat”?

```
candy["Kit Kat", "winpercent"]
```

[1] 76.7686

Q5. What is the winpercent value for “Tootsie Roll Snack Bars”?

```
candy["Tootsie Roll Snack Bars", "winpercent"]
```

[1] 49.6535

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?

Yes! Winpercent seems to be out of 100, which is very different from other columns.

Q7. What do you think a zero and one represent for the candy\$chocolate column?

candy\$chocolate

```
[1] 1 1 0 0 0 1 1 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 1 1 1 1 0 1 1 0 0 0 1 1 0 1 1 1  
[39] 1 1 1 0 1 1 0 0 0 1 0 0 0 1 1 1 1 0 1 0 0 1 0 0 1 0 1 1 0 0 0 0 0 0 0 0 1 1  
[77] 1 1 0 1 0 0 0 0 1
```

The 1 represents TRUE while the 0 represents FALSE.

Exploratory analysis

Q8. Plot a histogram of winpercent values.


```
hist(candy$winpercent)
```



```
library(ggplot2)

ggplot(candy) +
  aes(winpercent, bins=30) +
  geom_histogram()
```

`stat_bin()` using `bins = 30`. Pick better value `binwidth`.



Q9. Is the distribution of winpercent values symmetrical?

No, it is skewed to the left.

Q10. Is the center of the distribution above or below 50%?

```
mean(candy$winpercent)
```

[1] 50.31676

It is above 50%.

```
summary(candy$winpercent)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
22.45	39.14	47.83	50.32	59.86	84.18

Q11. On average is chocolate candy higher or lower ranked than fruit candy?

1. Find all chocolate candy
2. Get their winpercent values
3. Find the mean

4. Find all fruity candy
5. Get their winpercent values
6. Find the mean
7. Compare the two means

```
choc.candy <- candy[candy$chocolate ==1, ]
choc.win <- choc.candy$winpercent
mean(choc.win)
```

[1] 60.92153

```
fruity.candy <- candy[candy$fruity ==1, ]
fruity.win <- fruity.candy$winpercent
mean(fruity.win)
```

[1] 44.11974

Chocolate is higher ranked than fruity candy.

Q12. Is this difference statistically significant?

```
t.test(choc.win, fruity.win)
```

Welch Two Sample t-test

```
data: choc.win and fruity.win
t = 6.2582, df = 68.882, p-value = 2.871e-08
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
11.44563 22.15795
sample estimates:
mean of x mean of y
60.92153 44.11974
```

Yes, they are statistically different cause p-value is less than 0.05.

Overall Candy Rankings

Q13. What are the five least liked candy types in this set?

```
y <- c("y", "a", "z")
sort(y)
```

```
[1] "a" "y" "z"
```

```
y
```

```
[1] "y" "a" "z"
```

```
order(y)
```

```
[1] 2 1 3
```

```
ord.ind <- order(candy$winpercent)
head(candy[ord.ind, ], 5)
```

	chocolate	fruity	caramel	peanut	yalmond	nougat
Nik L Nip	0	1	0		0	0
Boston Baked Beans	0	0	0		1	0
Chiclets	0	1	0		0	0
Super Bubble	0	1	0		0	0
Jawbusters	0	1	0		0	0
	crispedrice	wafers	hard	bar	pluribus	sugarpercent
Nik L Nip	0	0	0	1	0.197	0.976
Boston Baked Beans	0	0	0	1	0.313	0.511
Chiclets	0	0	0	1	0.046	0.325
Super Bubble	0	0	0	0	0.162	0.116
Jawbusters	0	1	0	1	0.093	0.511
	winpercent					
Nik L Nip	22.44534					
Boston Baked Beans	23.41782					
Chiclets	24.52499					
Super Bubble	27.30386					
Jawbusters	28.12744					


Q14. What are the top 5 all time favorite candy types out of this set?

```
tail(candy[ord.ind, ], 5)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat
Snickers	1	0	1		1	1
Kit Kat	1	0	0		0	0
Twix	1	0	1		0	0
Reese's Miniatures	1	0	0		1	0
Reese's Peanut Butter cup	1	0	0		1	0
	crispedrice	wafer	hard bar	pluribus	sugar	percent
Snickers	0	0	1		0	0.546
Kit Kat	1	0	1		0	0.313
Twix	1	0	1		0	0.546
Reese's Miniatures	0	0	0		0	0.034
Reese's Peanut Butter cup	0	0	0		0	0.720
	price	percent	win	percent		
Snickers	0.651		76.67378			
Kit Kat	0.511		76.76860			
Twix	0.906		81.64291			
Reese's Miniatures	0.279		81.86626			
Reese's Peanut Butter cup	0.651		84.18029			


Q15. Make a first barplot of candy ranking based on winpercent values.

```
ggplot(candy) +  
  aes(winpercent, rownames(candy)) +  
  geom_col()
```




Q16. This is quite ugly, use the reorder() function to get the bars sorted by winpercent?

```
ggplot(candy) +
  aes(winpercent, reorder(rownames(candy),winpercent)) +
  geom_col()
```



```
ggplot(candy) +
  aes(winpercent,
      reorder(row.names(candy), winpercent),
      fill=chocolate)+  
  geom_col()
```



```
ylab("")
```

```
<ggplot2::labels> List of 1
$ y: chr ""
```


We need a custom color vector

```
my_cols <- rep("black", nrow(candy))
my_cols[candy$chocolate==1] <- "chocolate"
my_cols[candy$bar==1] <- "brown"
my_cols[candy$fruity==1] <- "darkgreen"
my_cols
```

```
[1] "brown"      "brown"       "black"       "black"       "darkgreen"   "brown"
[7] "brown"      "black"       "black"       "darkgreen"   "brown"       "darkgreen"
[13] "darkgreen"  "darkgreen"   "darkgreen"   "darkgreen"   "darkgreen"   "darkgreen"
[19] "darkgreen"  "black"       "darkgreen"   "darkgreen"   "chocolate"   "brown"
[25] "brown"      "brown"       "darkgreen"   "chocolate"   "brown"       "darkgreen"
[31] "darkgreen"  "darkgreen"   "chocolate"   "chocolate"   "darkgreen"   "chocolate"
[37] "brown"      "brown"       "brown"       "brown"       "brown"       "darkgreen"
[43] "brown"      "brown"       "darkgreen"   "darkgreen"   "brown"       "chocolate"
```

```
[49] "black"      "darkgreen" "darkgreen" "chocolate" "chocolate" "chocolate"
[55] "chocolate"  "darkgreen" "chocolate" "black"     "darkgreen" "chocolate"
[61] "darkgreen"  "darkgreen" "chocolate" "darkgreen" "brown"    "brown"
[67] "darkgreen"  "darkgreen" "darkgreen" "darkgreen" "black"    "black"
[73] "darkgreen"  "darkgreen" "darkgreen" "chocolate" "chocolate" "brown"
[79] "darkgreen"  "brown"    "darkgreen" "darkgreen" "darkgreen" "black"
[85] "chocolate"
```

```
ggplot(candy) +
  aes(winpercent,
      reorder(row.names(candy), winpercent),
      )+
  geom_col(fill=my_cols)
```



```
ylab("")
```

```
<ggplot2::labels> List of 1
$ y: chr ""
```

Q17. What is the worst ranked chocolate candy?

Sixlets

Q18. What is the best ranked fruity candy?


Starburst

Taking a look at pricepercent

```
library(ggrepel)
my_cols[candy$fruity==1] <- "red"

# How about a plot of win vs price
ggplot(candy) +
  aes(winpercent, pricepercent, label=rownames(candy)) +
  geom_point(col=my_cols) +
  geom_text_repel(col=my_cols, size=3.3, max.overlaps = 5)
```

Warning: ggrepel: 65 unlabeled data points (too many overlaps). Consider increasing max.overlaps



Q19. Which candy type is the highest ranked in terms of winpercent for the least money - i.e. offers the most bang for your buck?

Reese's Miniatures

Q20. What are the top 5 most expensive candy types in the dataset and of these which is the least popular?

```
ord <- order(candy$pricepercent, decreasing = TRUE)
head( candy[ord,c(11,12)], n=5 )
```

	pricepercent	winpercent
Nik L Nip	0.976	22.44534
Nestle Smarties	0.976	37.88719
Ring pop	0.965	35.29076
Hershey's Krackel	0.918	62.28448
Hershey's Milk Chocolate	0.918	56.49050

Nik L Nip, Ring pop, Nestle Smarties, Herhsey's Krackel

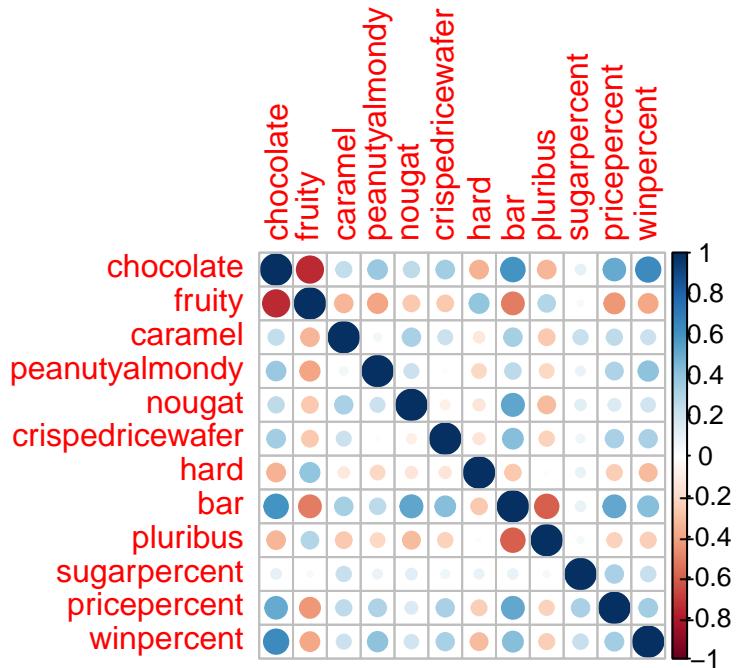
Exploring the correlation structure

```
cij <- cor(candy)
```

```
library(corrplot)
```

corrplot 0.95 loaded

```
corrplot(cij)
```



Q22. Examining this plot what two variables are anti-correlated (i.e. have minus values)?

chocolate and fruity

Q23. Similarly, what two variables are most positively correlated?

Chocolate and bar

Principal Component Analysis (PCA)

```
pca <- prcomp(candy, scale=TRUE)
summary(pca)
```

Importance of components:


	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369
	PC8	PC9	PC10	PC11	PC12		
Standard deviation	0.74530	0.67824	0.62349	0.43974	0.39760		
Proportion of Variance	0.04629	0.03833	0.03239	0.01611	0.01317		

```
Cumulative Proportion 0.89998 0.93832 0.97071 0.98683 1.00000
```


Score plot...

```
p <- ggplot(pca$x) +  
  aes(PC1, PC2, label=row.names(pca$x))+  
  geom_point(col=my_cols) +  
  geom_text_repel(max.overlaps=5, size=3.3, col=my_cols)  
  
p
```

Warning: ggrepel: 66 unlabeled data points (too many overlaps). Consider increasing max.overlaps



```
ggplot(pca$rotation) +  
  aes(PC1,  
       reorder(row.names(pca$rotation), PC1)) +  
  geom_col()
```



```
#library(plotly)
#ggplotly(p)
```

Q24. Complete the code to generate the loadings plot above. What original variables are picked up strongly by PC1 in the positive direction? Do these make sense to you? Where did you see this relationship highlighted previously?

Fruity, pluribus, and hard are picked up strongly in the positive direction. These make sense to me because they are characteristics of a fruity candy. These features/relationships were highlighted by the correlation plot previously.

Summary

Q25. Based on your exploratory analysis, correlation findings, and PCA results, what combination of characteristics appears to make a “winning” candy? How do these different analyses (visualization, correlation, PCA) support or complement each other in reaching this conclusion?

Combination of chocolate and bar seems to be the winning candy. This is seen with the correlation plot where these two characteristics combined had the deepest shade of blue coloring. The PCA also strongly picked up these combinations in the negative direction, stronger than any fruity candies. In the visualization, there were also more brown coloration and more were in the winpercent area, indicating that chocolate is both the majority and popularity.